Estimating the number of COVID-19 victims by combining the general logistic model with the pCN Monte Carlo sampling

May 28, 2020

1 Introduction

A simply effective way to quantify an epidemic impact, is given by counting the number of deaths. We are interested in using that number to better understand in which scenario we are likely to fall in the short and long term future.

To model the number of victims in time, we choose a generalized logistic map $X^P(t)$ governed by an unknown parameter $P \in \mathbb{R}^3$. Then, by using the Preconditioned Crank-Nicolson MCMC method on the available data, we establish what are the various possibilities for the values of P. Finally, by using the just learned values, running the model with them gives the various possibilities for the number of victims in the upcoming future.

Concretely speaking, we work with the most afflicted European Countries (Italy, Spain, Germany, France, UK). Our final goal is to answer to the following question: under the assumptions above, **how many n weeks of data do we need if we want to produce a prediction lasting for (at least) the next m weeks**?

We conclude by showing encouraging numerical results, suggesting that the chosen way might be good, but there are still possibilities of improvement.

2 The generalized logistic model

We choose to model the number of deaths in time with a generalized logistic function. Recall that these maps are all S-shaped curves, intuitively seen as an exponential start, then slowed down until reaching an horizontal equilibrium.

We do not provide any deep epidemiological explanation for such a choice, except that we expect a similar qualitative behavior for the phenomenon under analysis: the number of deaths is supposed to be higher at the beginning, reaching then a saturation point as long as the health system becomes capable of managing the emergency.

Of course, we are not new in proposing these maps: in the literature, there is a wide range of successful applications in similar fields (REF).

Our general logistic map is given by the following 1-dimension ODE:

$$X'(t) = \frac{q}{v} X(t) \left(1 - \left(\frac{X(t)}{Q}\right)^v\right) \tag{1}$$

with closed form solution:

$$X(t) = \frac{Q}{(1 + A\exp[-qt])^{\frac{1}{v}}} \tag{2}$$

for three (in theory) *stricly positivive* real parameters, $P = \{q, Q, v\}$ and initial time zero condition X_0 (we write $X(t) = X^P(t) = X^P_{X_0}(t)$ to stress this dependence). The letter A abbreviates $A = -1 + \left(\frac{Q}{X_0}\right)^v$.

Due to the limit $\lim_{t\to\infty} X(t) = Q$, we interpret Q as the maximum value reached asymptotically, sometimes referred as the *carrying capacity* in the literature.

The parameter v is related to the symmetry of the curve. Its limit for $v \to 0^+$ produces the *Gompertz* model, while the case $v = 1$ the *simple* logistic map.

The simple logistic model suggest an interpretation for the quantity q, too. In such a case the equation would just be

$$X'(t) = qX(t) \left(1 - \left(\frac{X(t)}{Q}\right)\right) \tag{3}$$

For very small time the value $X(t)$ is surely far lower than Q (otherwise the system would already be in equilibrium), therefore we would have $X'(t) \approx qX(t)$, giving so q as the rate of the starting exponential growth.

3 The bayesian approach

Let $X^P_{X_0}(t)$ be always the generalized logistic model above, depending on parameters P and with X_0 time zero condition. When using the model in practice, we can only observe a limited amount of points coming from its ODE trajectory, which are furthermore perturbed by a noise due to measurement errors. Let's fix $T + 1$ times $\{t_i\}_{i=0,\dots,T}$. In our experiments, T is generally between 7 and 21 days.

Definition 1. *The* observed vector $\boldsymbol{y} \in \mathbb{R}^{T+1}$ *is the random variable defined componentwise as:*

$$y_i(\omega) = X^P_{X_0}(t_i) + \eta_i(\omega) \tag{4}$$

where $\eta_i \sim \mathcal{N}(0, \sigma_i^2)$. [1]

[1] Sometimes by an abuse of notation we use the symbol η_i to indicate its density function too, so writing $\eta_i(x)$ for $x \in \mathbb{R}$ refers to that.

In other words, y represents the actual measurements on which we assume the influence of random errors. An important assumption that we make, lies in the belief that the errors follow the Gaussian structure above. In principle, many other possibilities can be chosen. In practice, the choice above is pretty standard and seems to work reasonabily well. We consider the measurements to be "reliable up to a 10% of error", in the sense that the "true" value at time i is supposed to be likely in the interval $y_i \pm \frac{y_i}{10}$. Recalling that $\eta_i \sim N(0, \sigma_i^2) \implies \mathbb{P}[-2\sigma_i \leq \eta_i \leq 2\sigma_i] \geq 95\%$, choosing then the standard deviation $\sigma_i \doteq \frac{y_i}{20}$ gives an error on y_i as desired.

Intuitively speaking, fitting the model means to estimate the unknown $P \in \mathbb{R}^3$ *given* the observation of $\mathbf{y}$. If we interpret P as a random variable, it means that the unknown quantity is given by the conditioned probability $\mathbb{P}[P|y]$.

We start by choosing a *prior* probability measure $\mathbb{P}[P]$ on $\mathbb{R}^3$, representing our blind guess about P *independently* of the observations $\mathbf{y}$. In practice this is sometimes an hard guess which can strongly influence the results. We will carefully describe our choice in a dedicated section.

If, as just pointed out, we aim at understanding $\mathbb{P}[P|\mathbf{y}]$, then the classical Bayes's law $\mathbb{P}[P|\mathbf{y}] \propto \mathbb{P}[\mathbf{y}|P] \times \mathbb{P}[P]$ allows to find it multiplying the the prior by an hypothetical conditioned law $\mathbb{P}[\mathbf{y}|P]$: this is where the notion of likelihood comes into play.

Definition 2. *For every fixed choice of the parameters P, the likelihood functions for the observation of $\boldsymbol{y}$, given P, is defined to be:*

$$\mathcal{L}(\boldsymbol{y}|P) \doteq \frac{(2\pi)^{-\frac{T}{2}}}{\sigma_0 \cdots \sigma_T} \exp\left(-\frac{1}{2}\sum_{i=0}^{T} \frac{(y_i - X^P_{X_0}(t_i))^2}{\sigma_i^2}\right) \tag{5}$$

By *interpreting* the likelihood as an effective probability conditioning, writing informally "$\mathbb{P}[\mathbf{y}|P] = \mathcal{L}(\mathbf{y}|P)$", not only the formula is explained when looking at the noise distribution $\eta(\mathbf{y} - X^P)$, but the fitting problem has now a clean solution:

Definition 3. *The (Bayesian) answer to the problem "Finding the probability density of the parameters P given the observations $\boldsymbol{y}$" is given by the* posterior *distribution on $\mathbb{R}^3$ defined by:*

$$\mu(dx) \propto \mathcal{L}(\boldsymbol{y}|x)\rho(dx) \tag{6}$$

where $\mathcal{L}$ is the likelihood defined above and ρ the prior probability for P.

During every use of the Bayesian rule, we constantly omitted the denominator relying always on the proportionality "$\propto$". This is because such a value is always the probability normalization constant, a number that can be completely ignored in practice thanks to the use of suitable Monte Carlo techniques.

3.1 The Preconditioned Crank-Nicolson MCMC

In the previous section we revised the Bayesian algorithm as a tool to convert the problem of estimating the parameter $\mathbb{P} \in \mathbb{R}^3$, to the task of sampling from the

posterior probability μ on $\mathbb{R}^3$. Its statistical properties convey the information that we need: for instance, the mode can be read like "the most probably choice for P", its variance can suggest how the uncertainty is spread, and so on. In principle one can analyze the posterior analytically, since it is explicitely given by formula NUMBER. But usually it is too hard to be done with pen and paper. Our answer it to use numerical tools: first we produce a large amount of samples, then we analyze them statistically.

In order to produce a (single) sample from μ, we use the so-called Preconditioned Crank-Nicolson Markov Chain Monte Carlo (pCN). Briefly speaking, it is a variant of the classic Gaussian Random Walk, but mainly used for much more complicated cases where for instance one can *arbitrarily* refine the ODE trajectory, or when the parameters P belongs to an infinite dimensional Banach space, or simply in the context of PDEs inverse problems. None of them is our case, since we have just $P \in \mathbb{R}^3$ and only limited daily available data.

Therefore there is no solid a priori justification for this specific Monte Carlo strategy, and one is free to choose another sampling method. We plan to test and study the various performances, but in another work: since the final results here are reasonabily fine, we decided to keep the algorithm for now, leaving potential room for improvement.

The reader is invited in consulting REF for a general exhaustive description of the pCN MCMC. In our simpler cases, the algorithm follows.

Given two candidate parameters $u, v \in \mathbb{R}^3$ define the acceptance probability:

$$a(u, v) \doteq \min\{1, \frac{\mathcal{L}(\mathbf{y}|v)}{\mathcal{L}(\mathbf{y}|u)}\} \tag{7}$$

and set the *exploratory* parameter $0 < \beta_{pcn} < 1$ (see step 3). The pCN *requires* the prior distribution $\rho(dx)$ to be a *centered* Gaussian, and used it as proposal step. This might be, in principle, a strong limitation: more comments about the prior will fillow in a dedicated section.

To produce a **single sample** from μ, construct a chain $\{x_i\}_{i\in\mathbb{N}}$ as follows:

1. set $x_0 \in \mathbb{R}^n$ **arbitrarily**. Then, for each $k > 0$:
2. sample a point $R \in \mathbb{R}^3$ from the gaussian prior distribution $\rho(dx)$;
3. propose a candidate as $\hat{x}_k = \sqrt{(1-\beta^2)}x_{k-1} + \beta R$;
4. accept it (i.e. set $x_k = \hat{x}_k$) with probability $a(x_{k-1}, \hat{x}_k)$;
5. (accepted or not) repeat from 2;

To avoid overflow problems, it is suggested to work with logatithms in the acceptance formula. Define N to be the integer at which we (always) stop the chain, so that everytime that we start from a point x_0, the chain produces a sample x_N. The value for N must be chosen in a way to overtake the known *burning time* issue. We always set it to be $2^15 > 130000$.

Repeating multiple times the algorithm with *different* starting points x_0, we collect large amounts of samples, call this number S, where the *correlation*

between them, a typical issue when when using a single traditional Markov Chain, is hopefully well mitigated. Again, in practice we set this value as 2^15, too.

Concerning the conservative parameter β, it was usually tuned around 0.001 or comparable values, chosen in order to produce a chain with an acceptance rate around 25%.

We still have to comment about the choice for the prior $\rho(dx)$ and the starting points x_0. Tt will be done in a section later.

We finally highlight how every chain istance is independent, therefore the complete sampling prodecure is suitable to parallelization (the user must be warned about the use of a proper seed).

4 Tuning the remaining parameters

4.1 Comments on the assumptions and methodology

We apply the previous theory in order to formulate hypothesis for multiple European Countries concerning the future number of deceased people. We firstly use part of the available data at the beginning of April to predict the behavior until the mid of May (now known), verify them, and then repeating the same methodology for the month of June. As we briefly said in the introduction, different number of weeks are taken as input data, in order to observe their predictions' strenght. A lot of careful points must be checked.

4.1.1 Why choosing an autonomous ODE

In the first section we justified the choice of a logistic map by looking at its S-shaped qualitative behavior, but we didn't comment the importance of keeping a curve described by an autonomous ODE.

Since the data in all Countries start around the mid of February, one might be tempted to use the *entire* collection (until, e.g. the 10th of April) relying on the intuition "more data, more accuracy".

As a first intuitive remark, recall that the general logistic model is connected to the simple logistic map ($v = 1$). For the latter one can show that, when used to model the number of infected, the map can be reformulated as a SIS compartmental model where q is also connected to the average number of people's daily social contacts. But since lockdown measures have been adopted, this quantity strongly changed in time. As a consequence of that, although we have no 100% rigorous justification, we prefer **not** to trust the entire available dataset, but rather limiting the observations on time period with *constant* (stronger or weaker) lockdown measures. Data for each Countries are given in their sections.

When following this logic, three cases can happen:

1 the "true" process is actually a general logistic one, but it already started *before* we began collecting data;

2 same as above, but the ODE *still* have to start, implying that part of the initial data that we read belongs to another model/trajectory;

3 the "true" model is not a general logistic one at all.

Of course, we cannot do nothing against 3, but the final results suggest that a general logistic model might be a "not-too-bad" choice. Conversely, by adding some days of delay we can increase the possibility of not being in 2. Finally, dealing with the problem number 1, it is necessary to ensure that the skipped days (e.g. including the one skypped in 2) can actually be "forgotten". But when choosing an autonomous ODE, this is precisely what claimed by the semigroup property of its associated flow. Briefly speaking, no matter if we start observing a value, say V_n, at day n, or V_{n+1} at day $n+1$: the trajectory produced from day $n+1$ (with initial conditions V_{n+1}) is precisely the same as the one produced by beginning at day n (with initial condition V_n). The two trajectories *are ruled by the same values of* P, therefore they infere the same parameter P (i.e. one can safely start the observations from the day $n+1$). Of course, one should not abuse this fact: the more are the observed data belonging the *same* trajectory, the better the estimation works. In other words, if one is sure that the ODE started at day n, one should take it as starting day and not $n+1$, in order to improve the estimator's performance.

4.2 Choosing the prior distribution

The goal of this section is to explain in which way we chosen the prior probability measure $\rho(dx)$. Since we used the pCN MCMC, it is required to be a centered Gaussian. It might represent a strong limitation, especially if we consider the fact that all the parameters in the general logistic equation are supposed to be positive. In order to cincumvent this problem, we tried with a very straightforward approach (which worked): we started the chain always with positive values, and when negative are proposed, we discard them.

Consequently we are technically not considering the full posterior distribution, but just its part with posisitive values. Altoight we still have to prove that there are no theoretical pitfalls, this naive approach seem to work reasonabily well.

Dealing with the precise structure of the prior distribution and the chain's starting points, everything is clarified in the following lines. Remember that the role of ρ is to represent the expected range of the searched values. For simplicity we define it with an easy covariance form:

$$\begin{pmatrix} \sigma_q^2 & 0 & 0 \\ 0 & \sigma_Q^2 & 0 \\ 0 & 0 & \sigma_v^2 \end{pmatrix} \tag{8}$$

splitting so the prior into three independent one dimensional Gaussians for every parameters. Recall that if $X \sim N(0, \sigma^2)$, the classical quantile property claims $\mathbb{P}[-2\sigma \leq X \leq 2\sigma] \geq 0.95$. In the first section we recall how we prefer having $v \in [0, 1]$, so by using the formula above we suggest $\sigma_v = 0.5$.

We remind that q is connected to the simple logistic's starting exponential growth, while Q relates to the maximum number of possible cases. Our idea is simple: as a first move, we do a rough exponential interpolation on the given dataset. Since we always obtained very small exponential powers, we use this knowledge to set $\sigma_q = 0.05$ (altought we are using a general logistic map, and not just the simple logistic model). Furthermore, the exponential law gives a rough prediction about the number of the deaths around, say, the $20th$ of May (i.e. one Month after the datasets); call it N_e. We set $\sigma_Q = \frac{N_e}{2}$, in order to represent the idea "a good candidate for the true number of deceased people is likely below the prediction done with the exponential case, being that the most aggressive accepted model".

Finally, it remains to tune the starting point $x_0 \in \mathbb{R}^3$ for every Markov Chain. Let D be the least read data from the dataset, i.e. the current number of deaths. Not surprisingly, we initialize x_0 uniformly randomly in $[0, 1] \times [D, N_e] \times [0, 1]$ coherently which the explanation above.

5 Numerical results

We took into account multiple European Countries with the idea of answering to the following question: supposing to be in a stable situation capable of avoiding an exponential explosion of the infection, how many weeks of data do we need before being able to perform realiable predictions concerning the number of deceased people? More precisely, for how many days do last the analysis done by using n weeks of data?

For each Country, we looked for the day when the lockdown measures came into play, waited around 15 days (incubation time, so to nullify the influence of the previous conditions), then we did predictions using 1, 2 and 3 weeks of data. They are generally done by using the days belonging to the first half of April, with the idea of lasting at least until the mid of May. Then all the Countries changed their policy. Therefore the idea is: if we manage to succed in predicting what happened in May by using the data in April, we can repat the same strategy concerning the month of June.

For each Country there is a dedicated section. Recall that the fitting procedure is done my using a Bayesian algorithm, producing so a probability measure for the parameters. We split it in clusters producing multi dimensional histograms, and considered only the 95% of its mass, eliminating so the very unlikely extreme value. We paid attention to the "worst" case scenario, i.e. the one with the highest number of deaths, the "best", and, when possible, the expectation value.

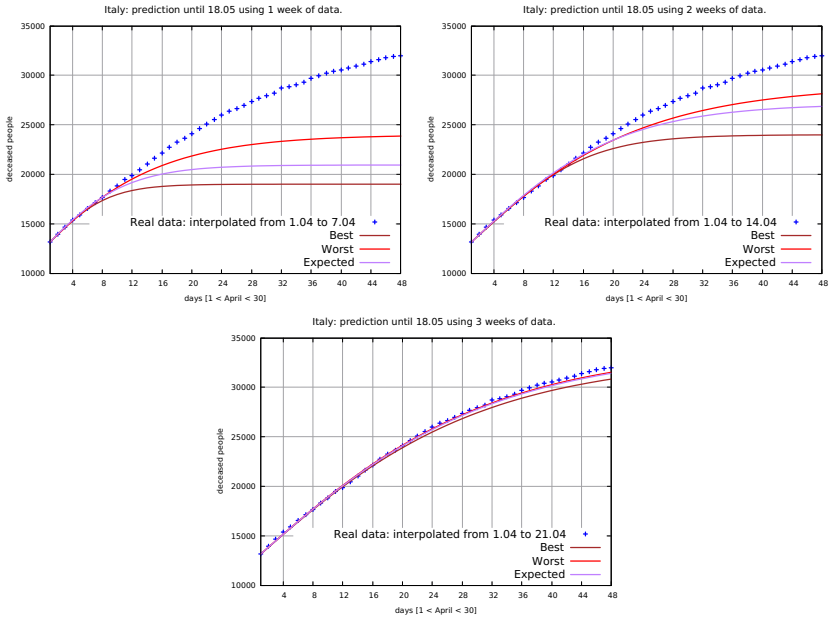

Figure 1: Italy - Lockdown measures were declared on the 15th of March. We started gathering the data from the 1st of April.

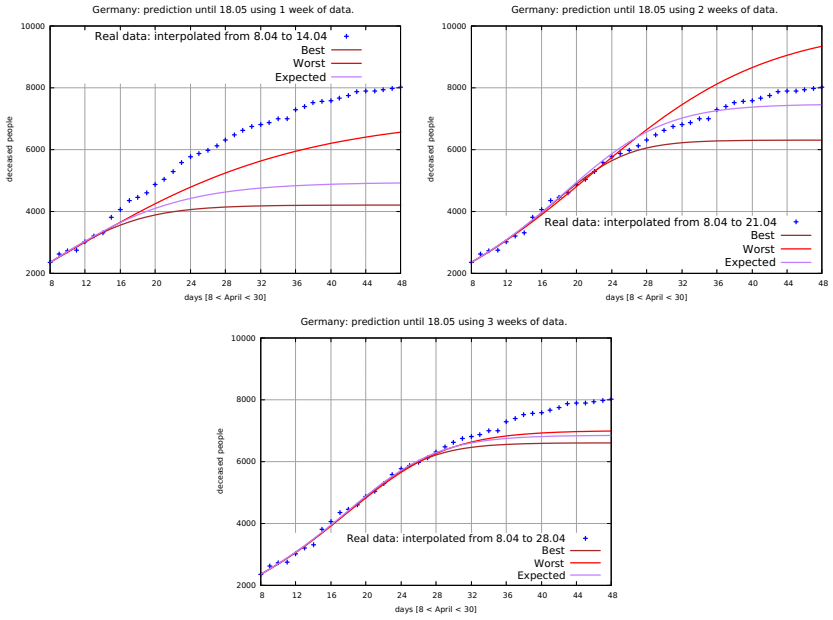

Figure 2: Germany -

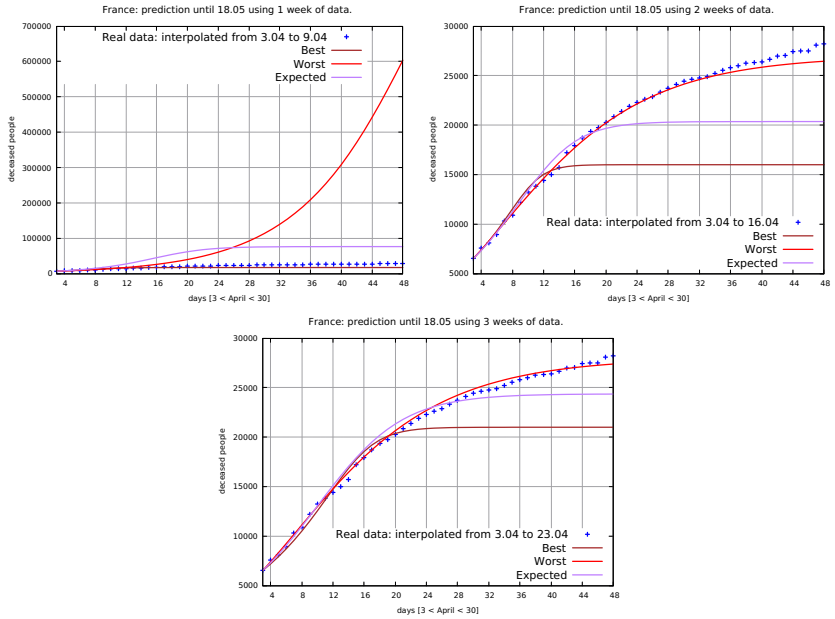

Figure 3: France - to write

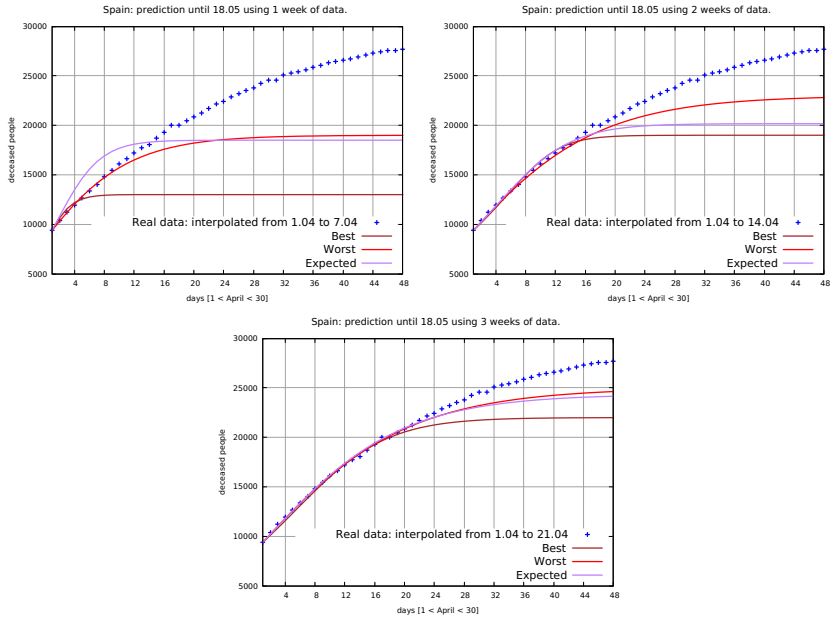

Figure 4: Spain - to write